

```
In [1]: #importing the libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: #Reading the dataset
df=pd.read_csv('day.csv')
df.head()
```

```
Out[2]:
```

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	temp
0	1	01-01-2018	1	0	1	0	1	1	2	14.1108
1	2	02-01-2018	1	0	1	0	2	1	2	14.9025
2	3	03-01-2018	1	0	1	0	3	1	1	8.0509
3	4	04-01-2018	1	0	1	0	4	1	1	8.2000
4	5	05-01-2018	1	0	1	0	5	1	1	9.3052

## 1. Inspecting the dataframe

```
In [3]: # Check the number of rows and columns in the dataframe
df.shape
```

```
Out[3]: (730, 16)
```

```
In [4]: # Check the column-wise info of the dataframe
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 730 entries, 0 to 729
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   instant    730 non-null    int64
 1   dteday     730 non-null    object
 2   season     730 non-null    int64
 3   yr         730 non-null    int64
 4   mnth       730 non-null    int64
 5   holiday    730 non-null    int64
 6   weekday    730 non-null    int64
 7   workingday 730 non-null    int64
 8   weathersit  730 non-null    int64
 9   temp       730 non-null    float64
10  atemp      730 non-null    float64
11  hum        730 non-null    float64
12  windspeed  730 non-null    float64
13  casual     730 non-null    int64
14  registered 730 non-null    int64
15  cnt        730 non-null    int64
dtypes: float64(4), int64(11), object(1)
memory usage: 91.4+ KB
```

```
In [5]: # Check the summary for the numeric columns
df.describe()
```

Out[5]:

	instant	season	yr	mnth	holiday	weekday	workingday
count	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000
mean	365.500000	2.498630	0.500000	6.526027	0.028767	2.995890	0.690411
std	210.877136	1.110184	0.500343	3.450215	0.167266	2.000339	0.462641
min	1.000000	1.000000	0.000000	1.000000	0.000000	0.000000	0.000000
25%	183.250000	2.000000	0.000000	4.000000	0.000000	1.000000	0.000000
50%	365.500000	3.000000	0.500000	7.000000	0.000000	3.000000	1.000000
75%	547.750000	3.000000	1.000000	10.000000	0.000000	5.000000	1.000000
max	730.000000	4.000000	1.000000	12.000000	1.000000	6.000000	1.000000

```
In [6]: #visualize missing values if any
import klib
klib.missingval_plot(df)
```

No missing values found in the dataset.

```
In [7]: # Converting date to Pandas datetime format
df['dteday'] = pd.to_datetime(df['dteday'])
```

In [8]: `df.head()`

Out[8]:

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	tenr
0	1	2018-01-01	1	0	1	0	1	1	2	14.11084
1	2	2018-02-01	1	0	1	0	2	1	2	14.90259
2	3	2018-03-01	1	0	1	0	3	1	1	8.05092
3	4	2018-04-01	1	0	1	0	4	1	1	8.20000
4	5	2018-05-01	1	0	1	0	5	1	1	9.30521

## 2. Data Cleaning and Analysis

In [9]: *#Changing the season, weathersit, mnth, weekday columns from numerical values to categorical strings*

```
df.season=df.season.map({1:'spring', 2:'summer', 3:'fall', 4:'winter'})
df.weathersit=df.weathersit.map({1:'Best', 2:'Neutral', 3:'Bad', 4:'Worse'})
df.mnth=df.mnth.map({1:'Jan',2:'Feb',3:'Mar',4:'Apr',5:'May',6:'June',7:'Jul',8:'Aug',9:'Sep',10:'Oct',11:'Nov',12:'Dec'})
df.weekday=df.weekday.map({1:'Mon',2:'Tue',3:'Wed',4:'Thu',5:'Fri',6:'Sat',0:'Sun'})
df.head()
```

Out[9]:

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	tenr
0	1	2018-01-01	spring	0	Jan	0	Mon	1	Neutral	14.11084
1	2	2018-02-01	spring	0	Jan	0	Tue	1	Neutral	14.90259
2	3	2018-03-01	spring	0	Jan	0	Wed	1	Best	8.05092
3	4	2018-04-01	spring	0	Jan	0	Thu	1	Best	8.20000
4	5	2018-05-01	spring	0	Jan	0	Fri	1	Best	9.30521

**Note: For the column weathersit, the alias is**

**Best :**

Clear, Few clouds, Partly cloudy, Partly cloudy

**Neutral:**

Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist

**Bad :**

Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rain + Scattered clouds

**Worse :**

Heavy Rain + Ice Pallets + Thunderstorm + Mist, Snow + Fog

```
In [10]: #The column 'instant' is very insignificant. Hence dropping that column.
df=df.drop('instant',axis=1)
df.shape
```

```
Out[10]: (730, 15)
```

```
In [11]: #Inserting a new variable day in the dataframe.
df.insert(4,'day','')
df['day']=pd.DatetimeIndex(df['dteday']).day
df.head()
```

```
Out[11]:
```

	dteday	season	yr	mnth	day	holiday	weekday	workingday	weathersit	temp
0	2018-01-01	spring	0	Jan	1	0	Mon	1	Neutral	14.110847
1	2018-02-01	spring	0	Jan	1	0	Tue	1	Neutral	14.902598
2	2018-03-01	spring	0	Jan	1	0	Wed	1	Best	8.050924
3	2018-04-01	spring	0	Jan	1	0	Thu	1	Best	8.200000
4	2018-05-01	spring	0	Jan	1	0	Fri	1	Best	9.305237

```
In [12]: print('-----day values-----')
print(df.day.value_counts())
print('-----workingday values-----')
print(df.workingday.value_counts())
print('-----weekday values-----')
print(df.weekday.value_counts())
```

```
-----day values-----
16    24
15    24
2     24
3     24
4     24
5     24
6     24
7     24
8     24
9     24
10    24
11    24
12    24
13    24
14    24
1     24
17    24
18    24
19    24
20    24
21    24
22    24
23    24
24    24
25    24
26    24
27    24
28    24
30    22
29    22
31    14
Name: day, dtype: int64
-----workingday values-----
1     504
0     226
Name: workingday, dtype: int64
-----weekday values-----
Mon     105
Tue     105
Wed     104
Fri     104
Sun     104
Thu     104
Sat     104
Name: weekday, dtype: int64
```

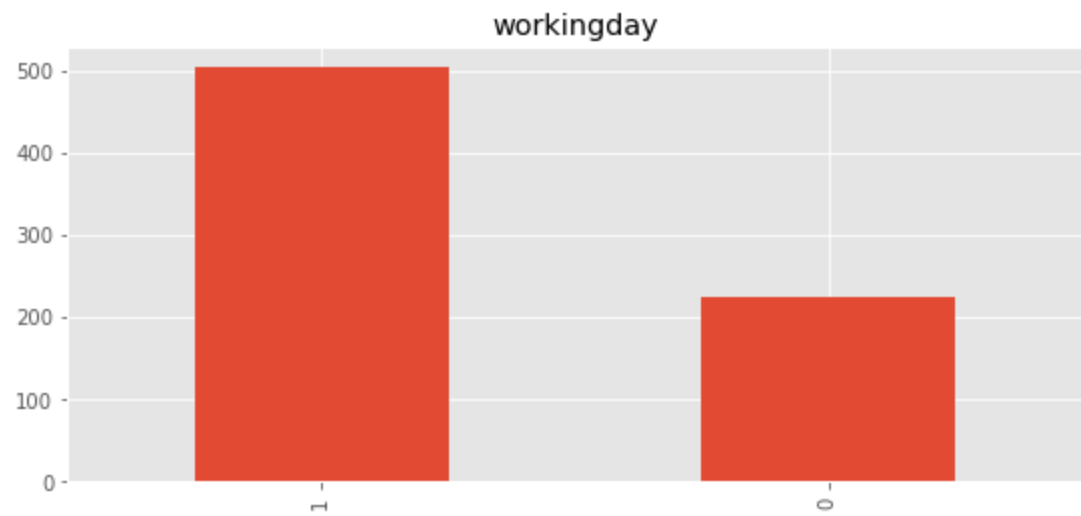
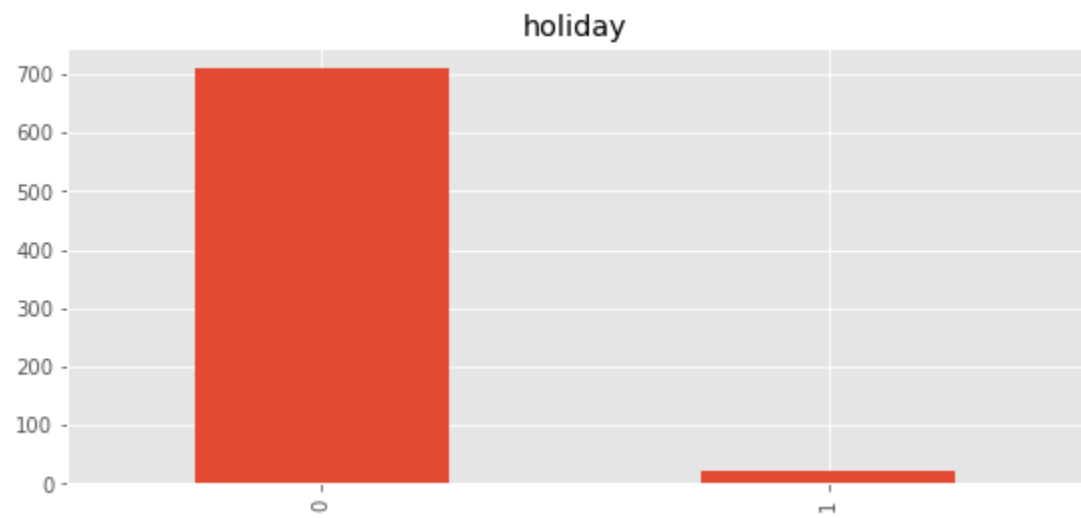

```
In [13]: #dropping dteday  
df=df.drop('dteday', axis=1)  
df.shape
```

```
Out[13]: (730, 15)
```

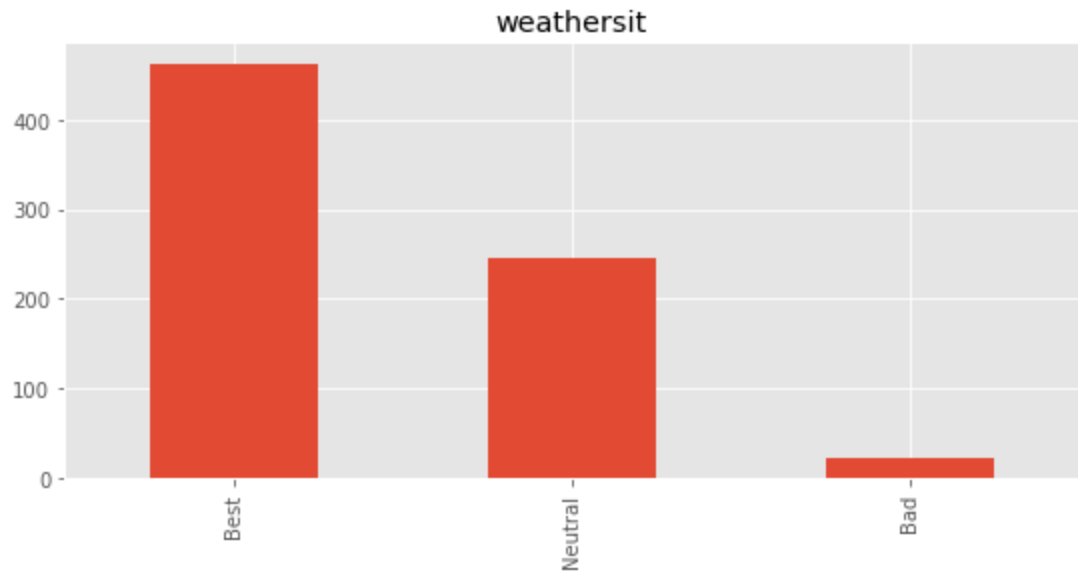
## Visualization

### 2.1 Univariate Analysis

```
In [14]: # Univariate analysis of few seemingly significant categorical variables:
univariate_categorical_cols=['season','holiday','workingday','weathersit']
plt.style.use('ggplot')
for column in univariate_categorical_cols:
    plt.figure(figsize=(10,4))
    plt.subplot(121)
    df[column].value_counts().plot(kind='bar')
    plt.title(column)
```







```
In [15]: print('Number of holidays in 2018: ', len(df[(df['holiday']==1) & (df['yr']==0)]))
          print('Number of holidays in 2019: ', len(df[(df['holiday']==1) & (df['yr']==1)]))
```

```
Number of holidays in 2018: 10
Number of holidays in 2019: 11
```

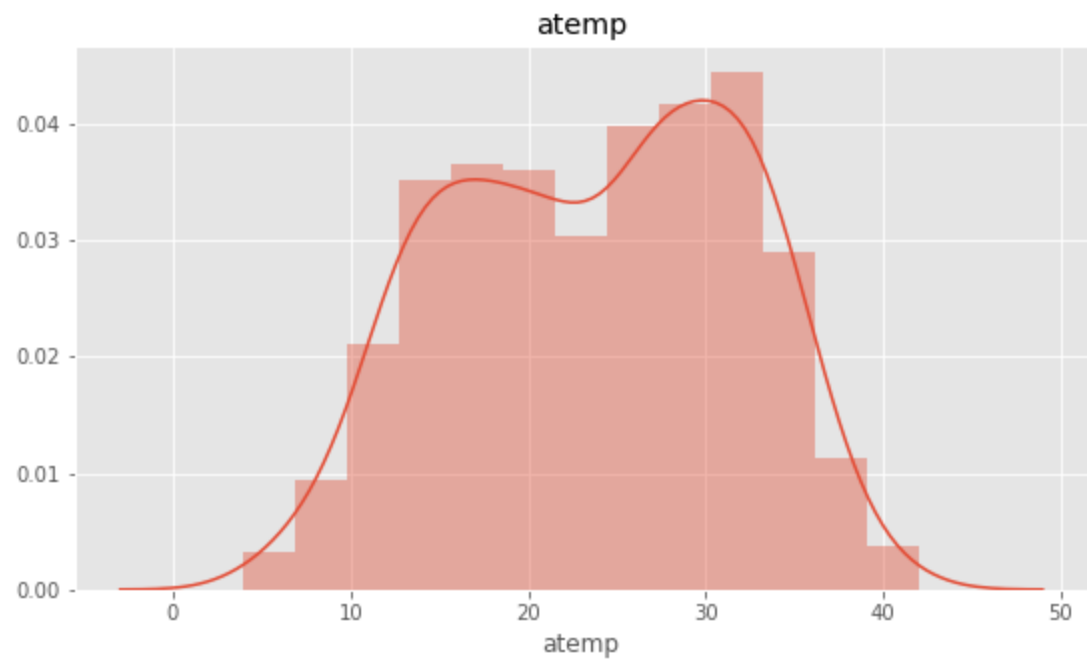

### Inferences:



1. Even though the margin is minimum, the number of days in fall is maximum and winter is minimum.  
Number of days as per season in decreasing order: Fall, Summer, Spring, Winter.
2. The number of public holidays is 21 in 2 years. Number of holidays in 2018 and 2019 are 10 and 11 respectively
3. The number of non-working days(Public holidays+weekends) is slightly less than half the number of working days which can be favourable for bike renting for exploring different places during non working days but can be non-favourable as well since the daily commute to office during the working days can be hampered.
4. Weather situation is mostly best case scenario and neutral compared to bad and worse which is favourable for renting bikes.



```
In [16]: # Univariate analysis of few seemingly significant continuous variables:
univariate_numerical_cols=df.select_dtypes(include=np.number)
univariate_numerical_cols=list(univariate_numerical_cols)
univariate_numerical_cols
univariate_continuous_var=[i for i in univariate_numerical_cols if
i not in ['yr',
'mnth',
'day',
'holiday',
'weekday',
'workingday']]
univariate_continuous_var
```


```
Out[16]: ['temp', 'atemp', 'hum', 'windspeed', 'casual', 'registered', 'cnt']
```

```
In [17]: for column in univariate_continuous_var:
          plt.figure(figsize=(20,5))
          plt.subplot(121)
          sns.distplot(df[column])
          plt.title(column)
```









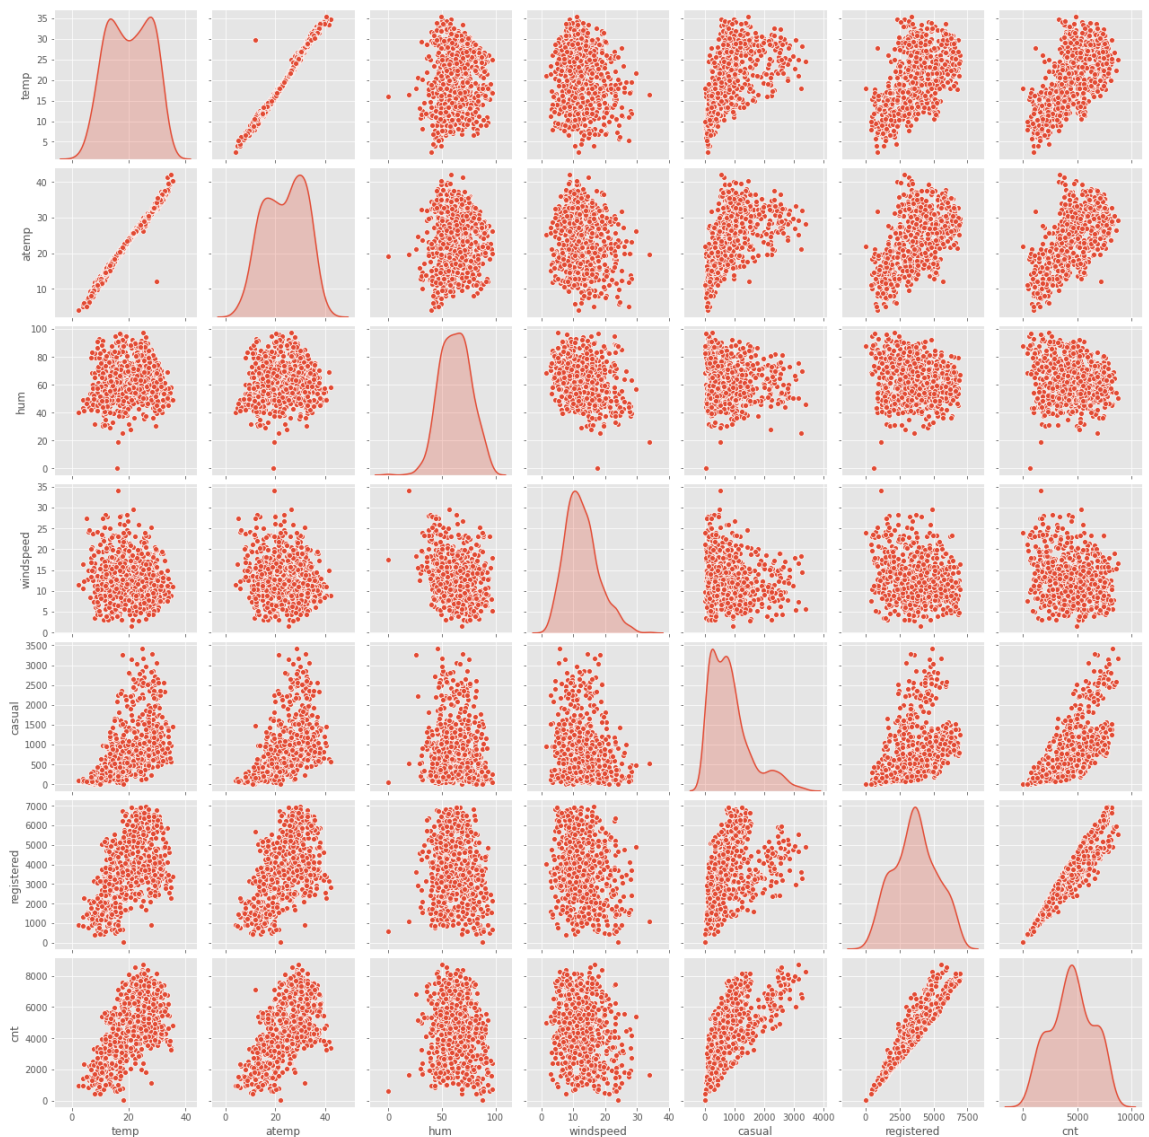
### Inferences:

1. Values of temperature and feeling temperature are differently distributed.
2. Humidity is almost randomly distributed with a mean of around 61-63.
3. The KDE of windspeed is almost a normal distribution with a right skew because of a few days with windspeed over 30.
4. The spread of casual users is not normally distributed whereas that of registered users is normally distributed ultimately leading to cnt to be spread normally distributed.

## 2.2 Bivariate Analysis

```
In [18]: df_continuous=df[univariate_continuous_var]
```

```
In [19]: #Bivariate analysis of continuos variables with cnt
sns.pairplot(df_continuous,diag_kind='kde')
plt.show()
```



## Inference

1. Huge corelation between temp and atemp. Hence only one of the 2 variables will be in the model.
2. temp/atemp shows some linear relationship with cnt.
3. hum and windspeed doesn't show much of a linear relationship with cnt.
4. Casual and registered shows linear relationship with cnt out of which the linear relationship shown by registered users is very significant.
5. Rest there are not any significant linear relationships.



```
In [20]: #Bivariate analysis of categorical variables with cnt
plt.figure(figsize=(30,48))
plt.subplot(8,2,1)
sns.boxplot(x='yr', y='cnt', data=df)

plt.subplot(8,2,2)
sns.barplot(x='season', y='cnt', data=df)

plt.subplot(8,2,3)
sns.boxplot(x='holiday', y='cnt', data=df)

plt.subplot(8,2,4)
sns.boxplot(x='weathersit', y='cnt', data=df)

plt.subplot(8,2,5)
sns.barplot(x='weathersit', y='windspeed', data=df)

plt.subplot(8,2,6)
sns.boxplot(x='workingday', y='cnt', data=df)

plt.subplot(8,2,7)
sns.barplot(x='mnth', y='windspeed', data=df)

plt.subplot(8,2,8)
sns.barplot(x='season', y='windspeed', data=df)

plt.subplot(8,2,9)
sns.lineplot(x='day', y='cnt', data=df)

plt.subplot(8,2,10)
sns.boxplot(x='mnth', y='cnt', data=df)

plt.subplot(8,2,11)
sns.barplot(x='mnth', y='cnt', data=df)

plt.subplot(8,2,12)
sns.barplot(x='weekday', y='cnt', data=df)

















plt.subplot(8,2,13)
sns.barplot(x='mnth', y='casual', data=df)

plt.subplot(8,2,14)
sns.barplot(x='weekday', y='casual', data=df)

plt.subplot(8,2,15)
sns.barplot(x='mnth', y='registered', data=df)

plt.subplot(8,2,16)
sns.barplot(x='weekday', y='casual', data=df)

plt.show()
```




```
In [21]: print('-----Winter Months-----')
print('Months')
print(df[df['season']=='winter'].mnth.value_counts())
print('-----Spring Months-----')
print('Months')
print(df[df['season']=='spring'].mnth.value_counts())
print('-----Summer Months-----')
print('Months')
print(df[df['season']=='summer'].mnth.value_counts())
print('-----Fall Months-----')
print('Months')
print(df[df['season']=='fall'].mnth.value_counts())
```

```
-----Winter Months-----
Months
Oct      62
Nov      60
Dec      40
Sep      16
Name: mnth, dtype: int64
-----Spring Months-----
Months
Jan      62
Feb      56
Mar      40
Dec      22
Name: mnth, dtype: int64
-----Summer Months-----
Months
May      62
Apr      60
June     40
Mar      22
Name: mnth, dtype: int64
-----Fall Months-----
Months
Aug      62
Jul      62
Sep      44
June     20
Name: mnth, dtype: int64
```

## Inferences

1. The cnt in the year 2019 was way more than that in 2018. The 75th percentile of the cnt in 2018 is almost equivalent to 25 percentile in 2019.
2. Number of bikes booked according to seasons in a decreasing order: Fall, Summer, Winter and Spring.
3. The trend of increasing use of bike starts from january(lowest) till June then stays almost the same till september and then starts dropping. There's a scope to increase the bike usage in the months from january till May and from October to december. The drop of bike usage from october till December might be explained by the winter season and less bike usage from January to April might be explained by higher windspeed.
4. Days of the week doesn't matter much. Almost similar number of bikes are rented same number of times everyday in a week but Monday and tuesday have relatively less bookings.
5. The average count of bikes rented is more during non-public holidays.
6. The average count of bikes rented when the weather is situation is 'Clear, Few clouds, Partly cloudy, Partly cloudy' or 'Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist' termed as 'Best' and 'Neutral' is much more compared to other situations termed as 'Bad' and 'Worse' - Wind speed during the bad weather situations is more than 'Best' and 'neutral' weather situations and hence more number of bikes are rented in such situations.
7. The line-graph trend shows that the count of bikes rented is least from 1st-4th day, peaks from 6th-10th day in a month and again dips till 13th day and kind of stays almost constant throughout the month.
8. There was a drop of bike rents by casual users in the winters and in the first 2 months of spring and less drop in the registered users. These people may be regular office going people or fitness enthusiasts.

```
In [22]: #Checking the colinearity amongst the variables
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(),annot=True,cmap="YlGnBu")
plt.show()
```



Since we have casual+registered=cnt and inferences are built from casual and registered records, let's drop them since these columns seem irrelevant for the model. Also it is a given that increasing casual or registered users both will be profitable factor for the business.

```
In [23]: df=df.drop(['casual', 'registered'],axis=1)
df.head()
```

Out[23]:

	season	yr	mnth	day	holiday	weekday	workingday	weathersit	temp	atemp	cnt
0	spring	0	Jan	1	0	Mon	1	Neutral	14.110847	18.18125	81
1	spring	0	Jan	1	0	Tue	1	Neutral	14.902598	17.68695	61
2	spring	0	Jan	1	0	Wed	1	Best	8.050924	9.47025	41
3	spring	0	Jan	1	0	Thu	1	Best	8.200000	10.60610	51
4	spring	0	Jan	1	0	Fri	1	Best	9.305237	11.46350	41

Also temp and atemp are very highly correlated and their respective colinearities with cnt are also same. Hence dropping atemp since feeling temperature can be relatively less accurate compared to temperature.

```
In [24]: df=df.drop('atemp',axis=1)
```

In [25]: `df.head()`

Out[25]:

	season	yr	mnth	day	holiday	weekday	workingday	weathersit	temp	hum	wi
0	spring	0	Jan	1	0	Mon	1	Neutral	14.110847	80.5833	10
1	spring	0	Jan	1	0	Tue	1	Neutral	14.902598	69.6087	16
2	spring	0	Jan	1	0	Wed	1	Best	8.050924	43.7273	16
3	spring	0	Jan	1	0	Thu	1	Best	8.200000	59.0435	16
4	spring	0	Jan	1	0	Fri	1	Best	9.305237	43.6957	12

### 3. Preparing data for modelling

In [26]: *#Creating Dummy variables*

```
def dummies(x,dataframe):
    temp = pd.get_dummies(dataframe[x], drop_first = True)
    dataframe = pd.concat([dataframe, temp], axis = 1)
    dataframe.drop([x], axis = 1, inplace = True)
    return dataframe
# Applying the function to the bikeSharing

df = dummies('season',df)
df = dummies('mnth',df)
df = dummies('weekday',df)
df = dummies('weathersit',df)
df.head()
```

Out[26]:

	yr	day	holiday	workingday	temp	hum	windspeed	cnt	spring	summer	...
0	0	1	0	1	14.110847	80.5833	10.749882	985	1	0	...
1	0	1	0	1	14.902598	69.6087	16.652113	801	1	0	...
2	0	1	0	1	8.050924	43.7273	16.636703	1349	1	0	...
3	0	1	0	1	8.200000	59.0435	10.739832	1562	1	0	...
4	0	1	0	1	9.305237	43.6957	12.522300	1600	1	0	...

5 rows × 30 columns

In [27]: `df.shape`

Out[27]: (730, 30)

In [28]: `df.describe()`

Out [28]:

	yr	day	holiday	workingday	temp	hum	windspeed
<b>count</b>	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000	730.000000
<b>mean</b>	0.500000	15.720548	0.028767	0.690411	20.319259	62.765175	12.763620
<b>std</b>	0.500343	8.802278	0.167266	0.462641	7.506729	14.237589	5.195841
<b>min</b>	0.000000	1.000000	0.000000	0.000000	2.424346	0.000000	1.500244
<b>25%</b>	0.000000	8.000000	0.000000	0.000000	13.811885	52.000000	9.041650
<b>50%</b>	0.500000	16.000000	0.000000	1.000000	20.465826	62.625000	12.125325
<b>75%</b>	1.000000	23.000000	0.000000	1.000000	26.880615	72.989575	15.625589
<b>max</b>	1.000000	31.000000	1.000000	1.000000	35.328347	97.250000	34.000021

8 rows × 30 columns

### 3.1 Splitting the data into test and train

```
In [29]: import sklearn
from sklearn.model_selection import train_test_split

df_train, df_test= train_test_split(df,train_size=0.7, random_state
=100)
print(df_train.shape)
print(df_test.shape)

(510, 30)
(220, 30)
```

### 3.2 Rescaling the features:

```
In [30]: #Rescaling even the target variables since a target variable with a
         #large spread of values, in turn, may result
         #in large error gradient values causing weight values to change dra-
         #matically, making the learning process unstable.

         from sklearn.preprocessing import MinMaxScaler
         scaler=MinMaxScaler()

         need_rescale=['temp','hum','windspeed','cnt']
         df_train[need_rescale]=scaler.fit_transform(df_train[need_rescale])

         df_train.describe()
```

Out[30]:

	yr	day	holiday	workingday	temp	hum	windspeed
count	510.000000	510.000000	510.000000	510.000000	510.000000	510.000000	510.000000
mean	0.507843	15.631373	0.025490	0.711765	0.537440	0.650480	0.320883
std	0.500429	8.852724	0.157763	0.453386	0.225858	0.145846	0.169803
min	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	8.000000	0.000000	0.000000	0.339853	0.538643	0.199179
50%	1.000000	16.000000	0.000000	1.000000	0.542596	0.653714	0.296763
75%	1.000000	23.000000	0.000000	1.000000	0.735215	0.754830	0.414447
max	1.000000	31.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 30 columns

### 3.3 Splitting train dataset into X and y

```
In [31]: y_train=df_train.pop('cnt')
         X_train=df_train
```

## 4. Model Building



In [32]: *#Since the total number of variables are 30, using RFE to calculate the best 15 variables to be used for model building*

```
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression

lm=LinearRegression()
lm.fit(X_train,y_train)

rfe=RFE(lm, 15)
rfe=rfe.fit(X_train,y_train)

list(zip(X_train.columns,rfe.support_,rfe.ranking_))
```

Out[32]:

```
[('yr', True, 1),
 ('day', False, 15),
 ('holiday', True, 1),
 ('workingday', False, 3),
 ('temp', True, 1),
 ('hum', True, 1),
 ('windspeed', True, 1),
 ('spring', True, 1),
 ('summer', True, 1),
 ('winter', True, 1),
 ('Aug', False, 9),
 ('Dec', True, 1),
 ('Feb', False, 2),
 ('Jan', True, 1),
 ('Jul', True, 1),
 ('June', False, 11),
 ('Mar', False, 14),
 ('May', False, 8),
 ('Nov', True, 1),
 ('Oct', False, 12),
 ('Sep', True, 1),
 ('Mon', False, 7),
 ('Sat', False, 4),
 ('Sun', False, 5),
 ('Thu', False, 13),
 ('Tue', False, 6),
 ('Wed', False, 10),
 ('Best', True, 1),
 ('Neutral', True, 1)]
```

In [33]:

```
col = X_train.columns[rfe.support_]
col
```

Out[33]:

```
Index(['yr', 'holiday', 'temp', 'hum', 'windspeed', 'spring', 'summer',
      'winter', 'Dec', 'Jan', 'Jul', 'Nov', 'Sep', 'Best', 'Neutral'],
      dtype='object')
```

```
In [34]: X_train.columns[~rfe.support_]
```

```
Out[34]: Index(['day', 'workingday', 'Aug', 'Feb', 'June', 'Mar', 'May', 'Oct',
              'Mon',
              'Sat', 'Sun', 'Thu', 'Tue', 'Wed'],
              dtype='object')
```

```
In [35]: X_train_rfe=X_train[X_train.columns[rfe.support_]]
X_train_rfe.head()
```

```
Out[35]:
```

	yr	holiday	temp	hum	windspeed	spring	summer	winter	Dec	Jan	Jul	No
576	1	0	0.815169	0.725633	0.264686	0	0	0	0	0	1	
426	1	0	0.442393	0.640189	0.255342	1	0	0	0	0	0	
728	1	0	0.245101	0.498067	0.663106	1	0	0	1	0	0	
482	1	0	0.395666	0.504508	0.188475	0	1	0	0	0	0	
111	0	0	0.345824	0.751824	0.380981	0	1	0	0	0	0	

After minimizing the number of variables using RFE, using statsmodel to build an optimized model.

```
In [36]: #Defining 2 functions model and VIF to train model and calculate VIF
         #repeatedly.
         import statsmodels.api as sm

         def model(X,y):
             X=sm.add_constant(X)
             lm_model=sm.OLS(y,X).fit()
             print(lm_model.summary())
             return X

         from statsmodels.stats.outliers_influence import variance_inflation_factor

         def VIF(X):
             vif=pd.DataFrame()
             vif['Features']=X.columns
             vif['VIF']=[variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
             vif['VIF']=round(vif['VIF'],2)
             vif=vif.sort_values(by='VIF', ascending=False)
             return vif
```

First Model:

```
In [37]: #Training the first model  
X_train1=model(X_train_rfe,y_train)
```

# OLS Regression Results

```

=====
Dep. Variable:          cnt    R-squared:
0.845
Model:                  OLS    Adj. R-squared:
0.840
Method:                 Least Squares    F-statistic:
179.4
Date:                   Thu, 28 Jan 2021    Prob (F-statistic):
8.15e-189
Time:                   22:03:09    Log-Likelihood:
514.19
No. Observations:      510    AIC:
-996.4
Df Residuals:          494    BIC:
-928.6
Df Model:              15
Covariance Type:      nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                0.0732      0.048      1.540      0.124      -0.020
0.167
yr                  0.2304      0.008     28.487      0.000       0.215
0.246
holiday             -0.0911      0.026     -3.557      0.000      -0.141
-0.041
temp                0.4815      0.037     13.005      0.000       0.409
0.554
hum                -0.1622      0.038     -4.291      0.000      -0.236
-0.088
windspeed          -0.1887      0.026     -7.315      0.000      -0.239
-0.138
spring             -0.0613      0.021     -2.881      0.004      -0.103
-0.019
summer              0.0423      0.015      2.761      0.006       0.012
0.072
winter              0.1019      0.018      5.656      0.000       0.067
0.137
Dec                -0.0355      0.018     -2.024      0.043      -0.070
-0.001
Jan                -0.0434      0.018     -2.393      0.017      -0.079
-0.008
Jul                -0.0553      0.018     -3.030      0.003      -0.091
-0.019
Nov                -0.0387      0.019     -2.057      0.040      -0.076
-0.002
Sep                0.0755      0.017      4.466      0.000       0.042
0.109
Best                0.2465      0.026      9.331      0.000       0.195
0.298
Neutral            0.1922      0.025      7.687      0.000       0.143

```

0.241

```

=====
=====
Omnibus:                    66.656    Durbin-Watson:
2.025
Prob(Omnibus):              0.000    Jarque-Bera (JB):
161.040
Skew:                       -0.682    Prob(JB):
1.07e-35
Kurtosis:                   5.392    Cond. No.
26.0
=====
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [38]: *#Calculating Variance Inflation Factor*  
VIF(X\_train1)

Out [38]:

	Features	VIF
0	const	143.15
14	Best	10.36
15	Neutral	8.95
6	spring	5.27
3	temp	4.42
8	winter	3.83
7	summer	2.77
4	hum	1.92
12	Nov	1.77
10	Jan	1.68
9	Dec	1.50
11	Jul	1.49
13	Sep	1.34
5	windspeed	1.21
1	yr	1.04
2	holiday	1.03

```
In [39]: #VIF of Best > 10. But according to experience it seems people are  
         # more likely to use bikes in the best weather situations  
         # and hence seems significant.  
         # Let's drop Dec to see the difference in the significance of other  
         # variables and R squared  
         X_train1=X_train1.drop('Dec',axis=1)
```

## Second Model:

```
In [40]: X_train1=model(X_train1,y_train)
```

# OLS Regression Results

```

=====
Dep. Variable:          cnt    R-squared:
0.844
Model:                  OLS    Adj. R-squared:
0.839
Method:                 Least Squares    F-statistic:
190.8
Date:                  Thu, 28 Jan 2021    Prob (F-statistic):
4.41e-189
Time:                  22:03:09    Log-Likelihood:
512.08
No. Observations:      510    AIC:
-994.2
Df Residuals:          495    BIC:
-930.6
Df Model:              14
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                0.0629      0.047      1.326      0.185      -0.030
0.156
yr                   0.2302      0.008     28.371      0.000       0.214
0.246
holiday             -0.0920      0.026     -3.582      0.000      -0.142
-0.042
temp                0.5055      0.035     14.369      0.000       0.436
0.575
hum                -0.1697      0.038     -4.497      0.000      -0.244
-0.096
windspeed          -0.1858      0.026     -7.190      0.000      -0.237
-0.135
spring             -0.0562      0.021     -2.652      0.008      -0.098
-0.015
summer              0.0479      0.015      3.168      0.002       0.018
0.078
winter              0.0972      0.018      5.421      0.000       0.062
0.132
Jan                -0.0341      0.018     -1.936      0.053      -0.069
0.001
Jul                -0.0559      0.018     -3.057      0.002      -0.092
-0.020
Nov                -0.0236      0.017     -1.362      0.174      -0.058
0.010
Sep                 0.0802      0.017      4.775      0.000       0.047
0.113
Best                0.2404      0.026      9.131      0.000       0.189
0.292
Neutral             0.1876      0.025      7.511      0.000       0.139
0.237
=====

```



```

=====
Omnibus:                    60.634    Durbin-Watson:
2.047
Prob(Omnibus):              0.000    Jarque-Bera (JB):
138.746
Skew:                       -0.640    Prob(JB):
7.44e-31
Kurtosis:                   5.211    Cond. No.
25.9
=====
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [41]: VIF(X\_train1)

Out[41]:

	Features	VIF
0	const	141.50
13	Best	10.23
14	Neutral	8.88
6	spring	5.20
3	temp	3.97
8	winter	3.76
7	summer	2.68
4	hum	1.90
9	Jan	1.57
10	Jul	1.49
11	Nov	1.49
12	Sep	1.31
5	windspeed	1.21
1	yr	1.04
2	holiday	1.03

In [42]: *#R squared remained almost the same. Variable Nov seems insignificant*

In [43]: X\_train1=X\_train1.drop('Nov',axis=1)

Third model

```
In [44]: X_train1=model(X_train1,y_train)
```

# OLS Regression Results

```

=====
Dep. Variable:          cnt    R-squared:
0.843
Model:                  OLS    Adj. R-squared:
0.839
Method:                 Least Squares    F-statistic:
205.0
Date:                   Thu, 28 Jan 2021    Prob (F-statistic):
7.59e-190
Time:                   22:03:09    Log-Likelihood:
511.13
No. Observations:      510    AIC:
-994.3
Df Residuals:          496    BIC:
-935.0
Df Model:               13
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                0.0572      0.047      1.210      0.227      -0.036
0.150
yr                   0.2301      0.008     28.339      0.000       0.214
0.246
holiday              -0.0963      0.026     -3.773      0.000      -0.146
-0.046
temp                 0.5124      0.035     14.706      0.000       0.444
0.581
hum                  -0.1681      0.038     -4.452      0.000      -0.242
-0.094
windspeed            -0.1874      0.026     -7.253      0.000      -0.238
-0.137
spring               -0.0519      0.021     -2.476      0.014      -0.093
-0.011
summer               0.0502      0.015      3.336      0.001       0.021
0.080
winter               0.0919      0.018      5.247      0.000       0.057
0.126
Jan                  -0.0333      0.018     -1.892      0.059      -0.068
0.001
Jul                  -0.0556      0.018     -3.039      0.003      -0.092
-0.020
Sep                  0.0827      0.017      4.951      0.000       0.050
0.116
Best                 0.2392      0.026      9.084      0.000       0.187
0.291
Neutral              0.1866      0.025      7.469      0.000       0.138
0.236
=====

```

```

=====
Omnibus:              58.633    Durbin-Watson:

```

```

2.057
Prob(Omnibus):          0.000   Jarque-Bera (JB):
131.919
Skew:                   -0.626   Prob(JB):
2.26e-29
Kurtosis:               5.154   Cond. No.
25.8
=====
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [45]: VIF(X\_train1)

Out[45]:

	Features	VIF
0	const	140.41
12	Best	10.21
13	Neutral	8.87
6	spring	5.08
3	temp	3.89
8	winter	3.59
7	summer	2.65
4	hum	1.90
9	Jan	1.57
10	Jul	1.49
11	Sep	1.30
5	windspeed	1.21
1	yr	1.04
2	holiday	1.02

In [46]: *#R squared remained almost the same. Variable Jan seems insignificant*

In [47]: X\_train1=X\_train1.drop('Jan',axis=1)

#### Fourth Model

```
In [48]: X_train1=model(X_train1,y_train)
```

# OLS Regression Results

```

=====
Dep. Variable:          cnt    R-squared:
0.842
Model:                  OLS    Adj. R-squared:
0.838
Method:                 Least Squares    F-statistic:
220.6
Date:                   Thu, 28 Jan 2021    Prob (F-statistic):
2.95e-190
Time:                   22:03:09    Log-Likelihood:
509.29
No. Observations:       510    AIC:
-992.6
Df Residuals:           497    BIC:
-937.5
Df Model:               12
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025
const	0.0478	0.047	1.015	0.311	-0.045
yr	0.2294	0.008	28.208	0.000	0.213
holiday	-0.0969	0.026	-3.787	0.000	-0.147
temp	0.5299	0.034	15.728	0.000	0.464
hum	-0.1726	0.038	-4.569	0.000	-0.247
windspeed	-0.1822	0.026	-7.074	0.000	-0.233
spring	-0.0564	0.021	-2.700	0.007	-0.097
summer	0.0531	0.015	3.536	0.000	0.024
winter	0.0976	0.017	5.643	0.000	0.064
Jul	-0.0572	0.018	-3.123	0.002	-0.093
Sep	0.0833	0.017	4.973	0.000	0.050
Best	0.2369	0.026	8.983	0.000	0.185
Neutral	0.1843	0.025	7.364	0.000	0.135

```

=====
Omnibus:                57.486    Durbin-Watson:
2.051
Prob(Omnibus):           0.000    Jarque-Bera (JB):

```

```

130.221
Skew:                -0.612    Prob(JB):
5.28e-29
Kurtosis:            5.151    Cond. No.
25.7
=====
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [49]: VIF(X\_train1)

Out[49]:

	Features	VIF
0	const	138.87
11	Best	10.19
12	Neutral	8.85
6	spring	5.02
3	temp	3.61
8	winter	3.48
7	summer	2.62
4	hum	1.89
9	Jul	1.48
10	Sep	1.30
5	windspeed	1.19
1	yr	1.03
2	holiday	1.02

```
In [50]: #All the variables seems significant now after evaluating  $P > |t|$  and VIF

#R squared from model summary is 0.842
r2=0.842

#Calculating adjusted R squared:
n = X_train1.shape[0]

# Number of features (predictors, p) is the shape along axis 1
p = X_train1.shape[1]

# We find the Adjusted R-squared using the formula

adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)
adjusted_r2
```

Out [50]: 0.8378588709677419

The variables Best has a VIF slightly greater than 10. But best case weather scenario must be kept while building the model. Hence considering the above model to be the ideal one. But let's drop a few more variables to see the changes in R squared, F-statistic and Prob (F-statistic) and if we could come up with a better model

```
In [51]: X_train1=X_train1.drop('Best',axis=1)
```

### Fifth Model



```
In [52]: X_train1=model(X_train1,y_train)
```

# OLS Regression Results

```

=====
Dep. Variable:          cnt    R-squared:
0.816
Model:                  OLS    Adj. R-squared:
0.812
Method:                 Least Squares    F-statistic:
201.1
Date:                   Thu, 28 Jan 2021    Prob (F-statistic):
3.01e-175
Time:                   22:03:10    Log-Likelihood:
470.93
No. Observations:       510    AIC:
-917.9
Df Residuals:           498    BIC:
-867.0
Df Model:               11
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025
const	0.3419	0.037	9.366	0.000	0.270
yr	0.2299	0.009	26.257	0.000	0.213
holiday	-0.0869	0.028	-3.158	0.002	-0.141
temp	0.5685	0.036	15.795	0.000	0.498
hum	-0.3057	0.037	-8.169	0.000	-0.379
windspeed	-0.2292	0.027	-8.440	0.000	-0.283
spring	-0.0430	0.022	-1.915	0.056	-0.087
summer	0.0602	0.016	3.731	0.000	0.029
winter	0.1028	0.019	5.523	0.000	0.066
Jul	-0.0631	0.020	-3.196	0.001	-0.102
Sep	0.0815	0.018	4.519	0.000	0.046
Neutral	-0.0220	0.011	-2.062	0.040	-0.043

```

=====
Omnibus:                95.895    Durbin-Watson:
2.033
Prob(Omnibus):           0.000    Jarque-Bera (JB):
249.907
Skew:                    -0.933    Prob(JB):

```

```

5.41e-55
Kurtosis:                    5.877    Cond. No.
19.2

```

```

=====
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [53]: *#The value of R squared decreased and value of F-statistic dropped significantly which shows that the fourth model was more fit then the fifth. Still trying to drop spring to see if better model can be achieved.*

In [54]: VIF(X\_train1)

Out [54]:

	Features	VIF
0	const	71.84
6	spring	4.99
3	temp	3.56
8	winter	3.48
7	summer	2.61
4	hum	1.60
9	Jul	1.48
11	Neutral	1.39
10	Sep	1.30
5	windspeed	1.14
1	yr	1.03
2	holiday	1.01

In [55]: X\_train1=X\_train1.drop('spring',axis=1)

## Sixth Model

```
In [56]: X_train1=model(X_train1,y_train)
```

# OLS Regression Results

```

=====
Dep. Variable:          cnt    R-squared:
0.815
Model:                  OLS    Adj. R-squared:
0.811
Method:                 Least Squares    F-statistic:
219.7
Date:                  Thu, 28 Jan 2021    Prob (F-statistic):
1.22e-175
Time:                  22:03:10    Log-Likelihood:
469.06
No. Observations:      510    AIC:
-916.1
Df Residuals:          499    BIC:
-869.5
Df Model:              10
Covariance Type:      nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const                0.2961      0.028     10.704      0.000      0.242
0.350
yr                   0.2289      0.009     26.118      0.000      0.212
0.246
holiday              -0.0886      0.028     -3.213      0.001     -0.143
-0.034
temp                 0.6198      0.024     25.761      0.000      0.573
0.667
hum                  -0.3124      0.037     -8.361      0.000     -0.386
-0.239
windspeed            -0.2340      0.027     -8.631      0.000     -0.287
-0.181
summer               0.0819      0.012      7.096      0.000      0.059
0.105
winter               0.1312      0.011     11.562      0.000      0.109
0.153
Jul                  -0.0558      0.019     -2.872      0.004     -0.094
-0.018
Sep                  0.0914      0.017      5.279      0.000      0.057
0.125
Neutral              -0.0207      0.011     -1.933      0.054     -0.042
0.000
=====

```

```

=====
Omnibus:              93.118    Durbin-Watson:
2.045
Prob(Omnibus):        0.000    Jarque-Bera (JB):
227.239
Skew:                 -0.931    Prob(JB):
4.53e-50
Kurtosis:             5.689    Cond. No.

```

16.4

```
=====
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors
is correctly specified.
```

```
In [57]: #The value of R squared remained the same but the value of F-statistic
has increased and almost similar to our ideal fourth model.
#This can be our another ideal model.
#Even though Neutral is one more vital variable for our ideal fourth
model, its p-value is higher.
#Let's try to drop that variable.
```

```
In [58]: VIF(X_train1)
```

Out[58]:

	Features	VIF
0	const	41.04
4	hum	1.59
3	temp	1.58
8	Jul	1.43
10	Neutral	1.39
6	summer	1.33
7	winter	1.28
9	Sep	1.19
5	windspeed	1.13
1	yr	1.03
2	holiday	1.01

```
In [59]: X_train1=X_train1.drop('Neutral',axis=1)
```

## Seventh Model

```
In [60]: X_train1=model(X_train1,y_train)
```

# OLS Regression Results

=====					
=====					
Dep. Variable:	cnt	R-squared:			
0.814					
Model:	OLS	Adj. R-squared:			
0.810					
Method:	Least Squares	F-statistic:			
242.4					
Date:	Thu, 28 Jan 2021	Prob (F-statistic):			
4.86e-176					
Time:	22:03:10	Log-Likelihood:			
467.16					
No. Observations:	510	AIC:			
-914.3					
Df Residuals:	500	BIC:			
-872.0					
Df Model:	9				
Covariance Type:	nonrobust				
=====					
=====					
	coef	std err	t	P> t	[0.025
0.975]					
-----					
const	0.3098	0.027	11.552	0.000	0.257
0.362					
yr	0.2278	0.009	25.978	0.000	0.211
0.245					
holiday	-0.0868	0.028	-3.139	0.002	-0.141
-0.032					
temp	0.6283	0.024	26.480	0.000	0.582
0.675					
hum	-0.3492	0.032	-10.838	0.000	-0.412
-0.286					
windspeed	-0.2380	0.027	-8.778	0.000	-0.291
-0.185					
summer	0.0812	0.012	7.016	0.000	0.058
0.104					
winter	0.1334	0.011	11.788	0.000	0.111
0.156					
Jul	-0.0553	0.019	-2.841	0.005	-0.094
-0.017					
Sep	0.0910	0.017	5.240	0.000	0.057
0.125					
=====					
=====					
Omnibus:	87.662	Durbin-Watson:			
2.031					
Prob(Omnibus):	0.000	Jarque-Bera (JB):			
196.855					
Skew:	-0.909	Prob(JB):			
1.79e-43					
Kurtosis:	5.441	Cond. No.			
14.5					
=====					



=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [61]: VIF(X\_train1)

Out[61]:

	Features	VIF
0	const	38.36
3	temp	1.53
8	Jul	1.43
6	summer	1.33
7	winter	1.27
9	Sep	1.19
4	hum	1.18
5	windspeed	1.13
1	yr	1.03
2	holiday	1.01

This model again has lesser R squared than the fourth model but the F-statistic is much more than that.

**There are 2 models that can be considered as the best fits:**


**Fourth model and the Seventh model**

## 5. Residual Analysis of the trained data

```
In [62]: #Rebuilding the seventh model
lm_model7=sm.OLS(y_train,X_train1).fit()
y_train_pred7=lm_model7.predict(X_train1)
```

```
In [63]: %matplotlib inline
fig = plt.figure()
sns.distplot((y_train - y_train_pred7), bins = 20)
fig.suptitle('Error Terms', fontsize = 20)                                # Plot
heading
plt.xlabel('Errors', fontsize = 18)
```

Out[63]: Text(0.5, 0, 'Errors')



```
In [64]: #Rebuilding the fourth model
X_train_rfe=sm.add_constant(X_train_rfe)
X_train_rfe.head()
```

Out[64]:

	const	yr	holiday	temp	hum	windspeed	spring	summer	winter	Dec	Jan	
576	1.0	1	0	0.815169	0.725633	0.264686	0	0	0	0	0	
426	1.0	1	0	0.442393	0.640189	0.255342	1	0	0	0	0	
728	1.0	1	0	0.245101	0.498067	0.663106	1	0	0	1	0	
482	1.0	1	0	0.395666	0.504508	0.188475	0	1	0	0	0	
111	1.0	0	0	0.345824	0.751824	0.380981	0	1	0	0	0	

```
In [65]: X_train_rfe.drop(['Dec', 'Nov', 'Jan'], axis=1, inplace=True)
X_train_rfe.head()
```

Out[65]:


	const	yr	holiday	temp	hum	windspeed	spring	summer	winter	Jul	Sep	E
576	1.0	1	0	0.815169	0.725633	0.264686	0	0	0	1	0	
426	1.0	1	0	0.442393	0.640189	0.255342	1	0	0	0	0	
728	1.0	1	0	0.245101	0.498067	0.663106	1	0	0	0	0	
482	1.0	1	0	0.395666	0.504508	0.188475	0	1	0	0	0	
111	1.0	0	0	0.345824	0.751824	0.380981	0	1	0	0	0	

Columns of X\_train\_rfe are similar to the fourth model

```
In [66]: lm_model4=sm.OLS(y_train,X_train_rfe).fit()
y_train_pred4=lm_model4.predict(X_train_rfe)
```

```
In [67]: fig = plt.figure()
sns.distplot(y_train - y_train_pred4), bins = 20)
fig.suptitle('Error Terms', fontsize = 20) # Plot heading
plt.xlabel('Errors', fontsize = 18)
```

Out[67]: Text(0.5, 0, 'Errors')



Residual Analysis shows that error terms for both the models gives almost a normal distribution but the R squared value is better for the fourth model compared to the seventh model. Also normality of error distribution is slightly better for fourth model compared to seventh model.

Hence selecting the fourth model for prediction.

## 6. Making Predictions

### 6.1 Preparing data for prediction.

In [68]: `df_test.head()`

Out [68]:

	yr	day	holiday	workingday	temp	hum	windspeed	cnt	spring	summer	...
<b>184</b>	0	7	1	0	29.793347	63.7917	5.459106	6043	0	0	...
<b>535</b>	1	20	0	1	32.082500	59.2083	7.625404	6211	0	1	...
<b>299</b>	0	27	0	0	19.270000	81.2917	13.250121	2659	0	0	...
<b>221</b>	0	8	0	1	31.433347	42.4167	13.417286	4780	0	0	...
<b>152</b>	0	6	0	0	29.315000	30.5000	19.583229	4968	0	1	...

5 rows × 30 columns

```
In [69]: #rescaling columns from the list need_rescale=['temp','hum','windspeed','cnt']

df_test[need_rescale]=scaler.transform(df_test[need_rescale])

df_train.head()
```

Out [69]:

	yr	day	holiday	workingday	temp	hum	windspeed	spring	summer	winter	.
<b>576</b>	1	31	0	1	0.815169	0.725633	0.264686	0	0	0	.
<b>426</b>	1	3	0	0	0.442393	0.640189	0.255342	1	0	0	.
<b>728</b>	1	30	0	1	0.245101	0.498067	0.663106	1	0	0	.
<b>482</b>	1	28	0	0	0.395666	0.504508	0.188475	0	1	0	.
<b>111</b>	0	22	0	0	0.345824	0.751824	0.380981	0	1	0	.

5 rows × 29 columns

In [70]: `df_test.describe()`

Out[70]:

	yr	day	holiday	workingday	temp	hum	windspeed
<b>count</b>	220.000000	220.000000	220.000000	220.000000	220.000000	220.000000	220.000000
<b>mean</b>	0.481818	15.927273	0.036364	0.640909	0.558718	0.638221	0.313293
<b>std</b>	0.500809	8.700715	0.187620	0.480828	0.233187	0.148694	0.159584
<b>min</b>	0.000000	1.000000	0.000000	0.000000	0.046591	0.261915	-0.042808
<b>25%</b>	0.000000	9.000000	0.000000	0.000000	0.355429	0.529197	0.198843
<b>50%</b>	0.000000	15.500000	0.000000	1.000000	0.558172	0.625590	0.300126
<b>75%</b>	1.000000	24.000000	0.000000	1.000000	0.755981	0.743798	0.402718
<b>max</b>	1.000000	31.000000	1.000000	1.000000	0.984424	1.002146	0.807474

8 rows × 8 columns

## 6.2 Prediction with model 4

In [71]: `y_test=df_test.pop('cnt')`  
`X_train_rfe=X_train_rfe.drop('const',axis=1)`  
`X_test_model4=df_test[X_train_rfe.columns]`  
`X_test_model4.head()`

Out[71]:

	yr	holiday	temp	hum	windspeed	spring	summer	winter	Jul	Sep	Best	Ne
<b>184</b>	0	1	0.831783	0.657364	0.084219	0	0	0	1	0	0	
<b>535</b>	1	0	0.901354	0.610133	0.153728	0	1	0	0	0	1	
<b>299</b>	0	0	0.511964	0.837699	0.334206	0	0	1	0	0	0	
<b>221</b>	0	0	0.881625	0.437098	0.339570	0	0	0	0	0	1	
<b>152</b>	0	0	0.817246	0.314298	0.537414	0	1	0	0	0	1	

In [72]: `X_test_model4.shape`

Out[72]: (220, 12)

In [73]: `#Adding constant to dataframe`  
`X_test_model4=sm.add_constant(X_test_model4)`

In [74]: `#Prediction`  
`y_test_pred_model4=lm_model4.predict(X_test_model4)`

```
In [75]: #Calculating Test data R-squared:  
from sklearn.metrics import r2_score  
r2=r2_score(y_test, y_test_pred_model4)  
print(r2)
```

0.8151738700604121

```
In [76]: #Calculating adjusted R squared:  
n = X_test_model4.shape[0]  
  
# Number of features (predictors, p) is the shape along axis 1  
p = X_test_model4.shape[1]  
  
# Calculating Adjusted R-squared using the formula  
  
adjusted_r2 = 1-(1-r2)*(n-1)/(n-p-1)  
adjusted_r2
```

Out[76]: 0.8035100851613118


**Train R squared: 0.842**

**Train Adjusted R squared: 0.8378588709677419**

**Test R squared: 0.8378588709677419**

**Test Adjusted R squared: 0.8035100851613118**

```
In [77]: #Checking Homoscedasticity for train and test data
plt.figure(figsize=(12,8))
plt.subplot(1,2,1)
sns.scatterplot(y=y_train - y_train_pred4, x=y_train_pred4)
plt.subplot(1,2,2)
sns.scatterplot(y=y_test - y_test_pred_model4, x=y_test_pred_model4)
plt.show()
```



There is no clustering or pattern below or above 0.0 on the Y-axis. This model is giving best results compared to other 6 models formed earlier while training.


```
In [78]: # Evaluating the Algorithm
from sklearn import metrics
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_test_pred_model4))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_test_pred_model4))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred_model4)))
```

```
Mean Absolute Error: 0.0695347736271711
Mean Squared Error: 0.008837328237214432
Root Mean Squared Error: 0.094007064826078
```

Lower values of MAE, MSE and RMSE shows vouches for the good performance of the model.

```
In [79]: # understanding the spread.
fig = plt.figure()
sns.scatterplot(y_test, y_test_pred_model4)
fig.suptitle('y_test vs y_pred', fontsize=20)           # Plot heading
plt.xlabel('y_test', fontsize=18)                     # X-label
plt.ylabel('y_test_pred_model4', fontsize=16)
```

Out[79]: Text(0, 0.5, 'y\_test\_pred\_model4')



Based on the very close value of R squared and Adjusted R squared values of the train and test data sets and based on y\_test and y\_pred graph, it can be inferred that the our linear regression model has the below equation for it's best fitted line:

$$\text{cnt} = 0.0478 + 0.2294 \text{ yr} - 0.0969 \text{ holiday} + 0.5299 \text{ temp} - 0.1726 \text{ hum} - 0.1822 \text{ windspeed} - 0.0564 \text{ spring} + 0.0531 \text{ summer} + 0.0976 \text{ winter} - 0.0572 \text{ Jul} + 0.0833 \text{ sept} + 0.2369 (\text{Clear, Few clouds, Partly cloudy, Partly cloudy}) + 0.1843 (\text{Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist})$$



```
In [80]: print(lm_model4.summary())
```

# OLS Regression Results

```

=====
Dep. Variable:          cnt    R-squared:
0.842
Model:                  OLS    Adj. R-squared:
0.838
Method:                 Least Squares    F-statistic:
220.6
Date:                   Thu, 28 Jan 2021    Prob (F-statistic):
2.95e-190
Time:                   22:03:12    Log-Likelihood:
509.29
No. Observations:       510    AIC:
-992.6
Df Residuals:           497    BIC:
-937.5
Df Model:               12
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025
const	0.0478	0.047	1.015	0.311	-0.045
yr	0.2294	0.008	28.208	0.000	0.213
holiday	-0.0969	0.026	-3.787	0.000	-0.147
temp	0.5299	0.034	15.728	0.000	0.464
hum	-0.1726	0.038	-4.569	0.000	-0.247
windspeed	-0.1822	0.026	-7.074	0.000	-0.233
spring	-0.0564	0.021	-2.700	0.007	-0.097
summer	0.0531	0.015	3.536	0.000	0.024
winter	0.0976	0.017	5.643	0.000	0.064
Jul	-0.0572	0.018	-3.123	0.002	-0.093
Sep	0.0833	0.017	4.973	0.000	0.050
Best	0.2369	0.026	8.983	0.000	0.185
Neutral	0.1843	0.025	7.364	0.000	0.135

```

=====
Omnibus:                57.486    Durbin-Watson:
2.051
Prob(Omnibus):          0.000    Jarque-Bera (JB):

```

```

130.221
Skew:                -0.612    Prob(JB):
5.28e-29
Kurtosis:            5.151    Cond. No.
25.7
=====
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## Hypothesis Testing

Null hypothesis states that there is no relationship between the X variables and the Y variables meaning the coefficients of the independent variables is zero. From the final model summary, it is evident that all our coefficients are not equal to zero which means **We REJECT the NULL HYPOTHESIS**

The company should focus on the following factors:

1. People are less likely to use their service at low or extreme temperatures. So either the company can function to half the capacity or minimum capacity to reduce operational costs for better profits and provide service for regular registered customers mostly. Similarly in days with increase in humidity and windspeed. Discounts or offers won't help as well since it's inconvenient to commute using bikes in such situations.
2. There will be increase in the number of users with increase in year since people will start adapting to renting bikes more often. There might be chances that because of covid just been around the corner, the trend might not follow immediately but giving a year more will definitely see rise in number of users.
3. People are more likely to use their service in the best or the neutral weather environments i.e; Clear, Few clouds, Partly cloudy, Partly cloudy OR Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist.